

API Rate Limit Monitor - Complete Documentation

Table of Contents

1. [What is Rate Limiting?](#)
 2. [How This System Works](#)
 3. [Core Concepts](#)
 4. [Data Flow Examples](#)
 5. [Database Design](#)
 6. [API Endpoints Guide](#)
 7. [Real-World Scenario](#)
-

What is Rate Limiting?

Rate limiting is controlling how many requests a client can make to your API in a given time period.

Why Do We Need It?

Without Rate Limiting:

User A sends 10,000 requests per second → Server crashes
Malicious bot sends 1 million requests → Your bill explodes
Single user consumes all resources → Other users can't access the API

With Rate Limiting:

User A: Maximum 100 requests per minute → Controlled usage
Bot: Blocked after exceeding limit → System protected
Resources: Fairly distributed → All users get access

Real-World Examples

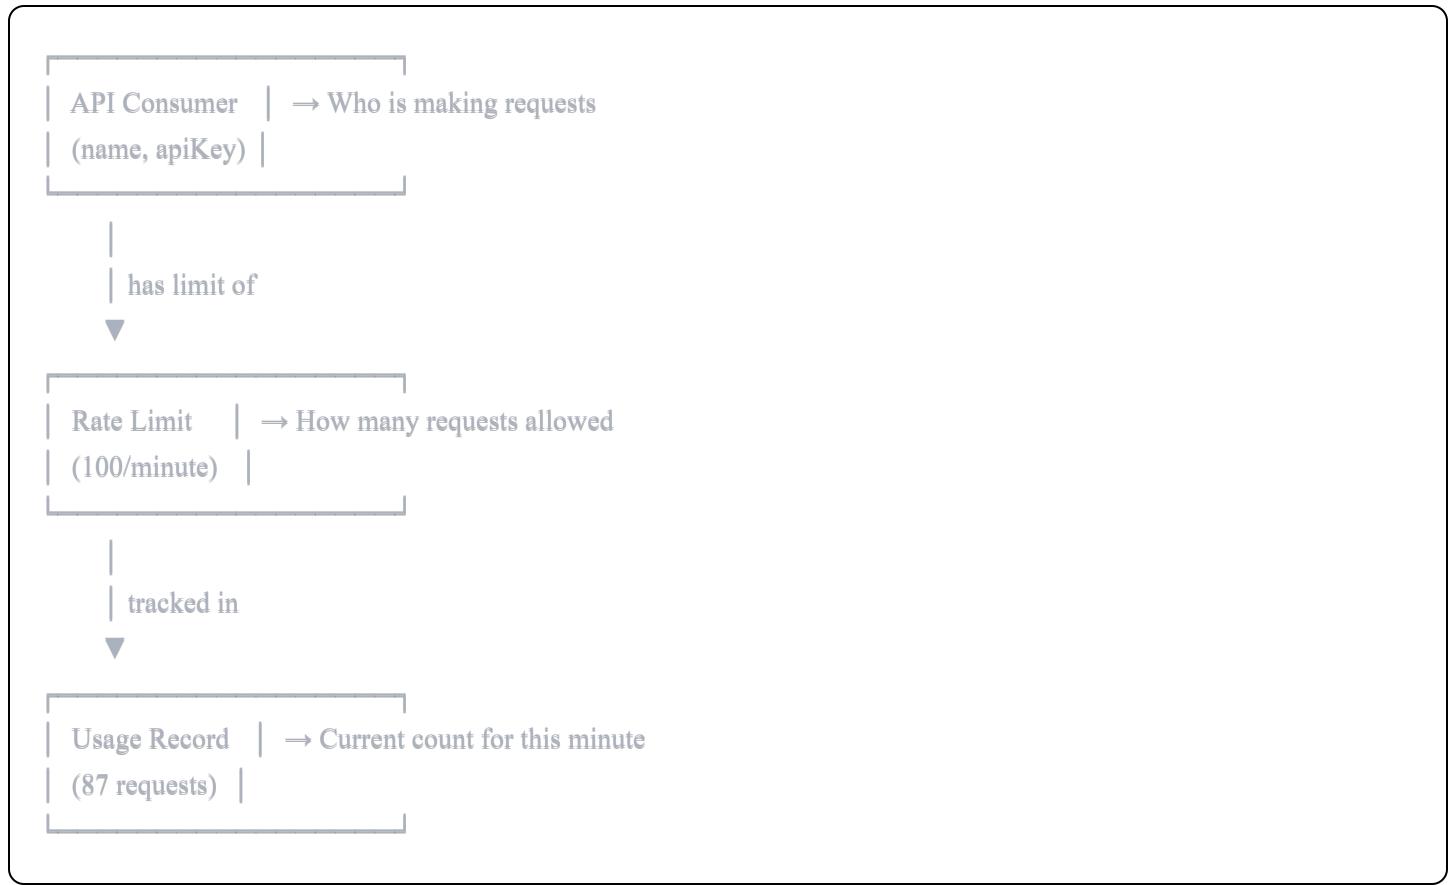
- **Twitter API:** 300 requests per 15 minutes
- **GitHub API:** 5,000 requests per hour
- **Google Maps API:** Based on your billing plan
- **Stripe API:** Different limits per endpoint

How This System Works

Think of it like a **ticket counter at a concert**:

1. **Registration:** Each person gets a ticket with their name (API Consumer with API Key)
2. **Entry Rules:** Each ticket allows 10 entries per hour (Rate Limit)
3. **Counter:** Someone tracks how many times you've entered (Usage Record)
4. **Time Windows:** The count resets every hour (Time Window)

System Components



Core Concepts

1. API Consumer

What is it? A user, application, or service that calls your API.

Example:

```
json
```

```
{  
  "id": 1,  
  "name": "Mobile App v2.0",  
  "apiKey": "a1b2c3d4e5f6",  
  "limitPerMinute": 100,  
  "status": "ACTIVE"  
}
```

Think of it as: A Netflix account. Each account has an ID, name, and subscription plan (rate limit).

2. API Key

What is it? A unique identifier that authenticates the consumer.

Example: `a1b2c3d4e5f6789012345678901234`

How it works:

Request without API Key → Rejected
Request with valid API Key → Proceed to rate limit check
Request with invalid API Key → Rejected (404 Not Found)

Think of it as: Your credit card number. Unique to you, used to identify you.

3. Time Windows

What is it? A fixed time period for counting requests.

Types:

- **MINUTE:** Resets every 60 seconds (10:30:00 → 10:31:00)
- **HOUR:** Resets every 60 minutes (10:00:00 → 11:00:00)

Example Timeline:



Think of it as: A gym membership that allows 10 visits per week. Every Monday, your counter resets to 0.

4. Usage Record

What is it? A record of how many requests a consumer made in a specific time window.

Database Row Example:

id	consumer_id	window_type	window_start	request_count
1	5	MINUTE	2026-01-18 10:30:00	87
2	5	MINUTE	2026-01-18 10:31:00	12
3	7	MINUTE	2026-01-18 10:30:00	5

Reading this table:

- Row 1: Consumer #5 made 87 requests between 10:30:00 and 10:31:00
- Row 2: Consumer #5 made 12 requests between 10:31:00 and 10:32:00
- Row 3: Consumer #7 made 5 requests between 10:30:00 and 10:31:00

Think of it as: A call log on your phone. It shows who called, when, and how many times.

Data Flow Examples

Example 1: Creating a New API Consumer

Step-by-Step:

1. User sends request:

```
POST /api/consumers
{
  "name": "Weather App",
  "limitPerMinute": 50
}
```

2. System generates unique API Key:

```
apiKey = "a1b2c3d4e5f6" (random UUID)
```

3. System saves to database:

```
INSERT INTO api_consumers (name, api_key, limit_per_minute, status)
VALUES ('Weather App', 'a1b2c3d4e5f6', 50, 'ACTIVE')
```

4. System responds:

```
{
  "id": 1,
  "name": "Weather App",
  "apiKey": "a1b2c3d4e5f6",
  "limitPerMinute": 50,
  "status": "ACTIVE"
}
```

Now the consumer can use **a1b2c3d4e5f6** to make API calls!

Example 2: First Request (Cold Start)

Scenario: Consumer makes their very first request at 10:30:15

1. Request arrives:

```
POST /api/rate-limit/record?apiKey=a1b2c3d4e5f6
```

2. System finds consumer:

```
SELECT * FROM api_consumers WHERE api_key = 'a1b2c3d4e5f6'
```

```
Result: {id: 1, limitPerMinute: 50, status: ACTIVE}
```

3. System checks if consumer is active:

Status is ACTIVE → Proceed

4. System calculates current time window:

Current time: 10:30:15

Window start: 10:30:00 (truncate to minute)

5. System looks for usage record:

```
SELECT * FROM usage_records
```

```
WHERE consumer_id = 1
```

```
AND window_type = 'MINUTE'
```

```
AND window_start = '2026-01-18 10:30:00'
```

```
Result: ✗ No record found (first request ever)
```

6. System creates new usage record:

```
INSERT INTO usage_records (consumer_id, window_type, window_start, request_count)
```

```
VALUES (1, 'MINUTE', '2026-01-18 10:30:00', 0)
```

7. System increments count:

```
UPDATE usage_records
```

```
SET request_count = 1
```

```
WHERE id = <new_id>
```

8. System checks limit:

$1 < 50 \rightarrow \checkmark$ ALLOWED

9. System responds:

```
{  
  "success": true,  
  "currentUsage": 1  
}
```

Database state after this request:

```
api_consumers:  
id | name      | api_key      | limit_per_minute | status  
1 | Weather App | a1b2c3d4e5f6 | 50          | ACTIVE
```

```
usage_records:  
id | consumer_id | window_type | window_start      | request_count  
1  | 1           | MINUTE       | 2026-01-18 10:30:00 | 1
```

Example 3: Multiple Requests in Same Minute

Scenario: Consumer makes 5 more requests at 10:30:20, 10:30:25, 10:30:40, 10:30:50, 10:30:55

Request #2 (10:30:20):

Current window: 10:30:00

Existing count: 1

New count: 2

Check: $2 < 50 \rightarrow \checkmark$ ALLOWED

Request #3 (10:30:25):

Current window: 10:30:00 (SAME WINDOW)

Existing count: 2

New count: 3

Check: $3 < 50 \rightarrow \checkmark$ ALLOWED

Request #4 (10:30:40):

Current window: 10:30:00 (SAME WINDOW)

Existing count: 3

New count: 4

Check: $4 < 50 \rightarrow \checkmark$ ALLOWED

... and so on

Database state:

```
usage_records:  
id | consumer_id | window_type | window_start      | request_count  
1  | 1           | MINUTE       | 2026-01-18 10:30:00 | 6
```

Key insight: The SAME row is updated because it's the same (consumer, window_type, window_start).

Example 4: New Minute Window

Scenario: Consumer makes a request at 10:31:05 (new minute!)

1. Request arrives at 10:31:05

2. System calculates window:

Current time: 10:31:05

Window start: 10:31:00 ← NEW WINDOW!

3. System looks for usage record:

```
SELECT * FROM usage_records  
WHERE consumer_id = 1  
AND window_type = 'MINUTE'  
AND window_start = '2026-01-18 10:31:00'  
Result: ❌ No record (different window!)
```

4. System creates NEW record:

```
INSERT INTO usage_records (consumer_id, window_type, window_start, request_count)  
VALUES (1, 'MINUTE', '2026-01-18 10:31:00', 1)
```

5. Response: ✅ ALLOWED (count reset to 1)

Database state:

usage_records:				
	id	consumer_id	window_type	window_start request_count
1	1	1	MINUTE	2026-01-18 10:30:00 6 ← OLD WINDOW
2	2	1	MINUTE	2026-01-18 10:31:00 1 ← NEW WINDOW

This is why the unique constraint is critical! It ensures we have exactly ONE row per time window.

Example 5: Rate Limit Exceeded

Scenario: Consumer has made 49 requests, now makes the 50th and 51st

Request #50:

Current count: 49

New count: 50

Check: $50 < 50 \rightarrow \text{X}$ FALSE, but we allow it ($49 + 1 = 50$ is at limit)

Actually, let's check the code...

Code says: if (requestCount \geq limitPerMinute) throw exception

So before recording:

$49 < 50 \rightarrow \checkmark$ Record it

Update count to 50

Request #51:

Current count: 50

Check: $50 \geq 50 \rightarrow \text{X}$ EXCEEDED!

Response: 429 TOO MANY REQUESTS

```
{  
  "error": "Rate limit exceeded"  
}
```

The user must wait until 10:31:00 for their counter to reset.

Database Design

Why Two Tables?

api_consumers = WHO can access the API **usage_records** = WHAT they've done recently

Separation of concerns:

- Consumer data changes rarely (name, limit)
- Usage data changes constantly (every request)

The Critical Constraint

sql

UNIQUE (consumer_id, window_type, window_start)

What it prevents:

X WITHOUT CONSTRAINT:

id	consumer_id	window_type	window_start	request_count
1	1	MINUTE	2026-01-18 10:30:00	25
2	1	MINUTE	2026-01-18 10:30:00	30 ← DUPLICATE!
3	1	MINUTE	2026-01-18 10:30:00	15 ← DUPLICATE!

Total requests: 70, but we can't know which is correct!

✓ WITH CONSTRAINT:

id	consumer_id	window_type	window_start	request_count
1	1	MINUTE	2026-01-18 10:30:00	70

Single source of truth! Accurate counting!

API Endpoints Guide

1. Create Consumer

Purpose: Register a new API consumer

Request:

```
http  
POST /api/consumers  
Content-Type: application/json  
  
{  
  "name": "My Mobile App",  
  "limitPerMinute": 100  
}
```

Response:

```
json
```

```
{  
  "id": 1,  
  "name": "My Mobile App",  
  "apiKey": "a1b2c3d4e5f6g7h8i9j0",  
  "limitPerMinute": 100,  
  "status": "ACTIVE"  
}
```

Use Case: When onboarding a new developer/application.

2. Check Rate Limit (Without Recording)

Purpose: Check if a consumer can make a request without actually counting it

Request:

```
http  
POST /api/rate-limit/check?apiKey=a1b2c3d4e5f6g7h8i9j0
```

Response (Under Limit):

```
json  
  
{  
  "allowed": true,  
  "currentUsage": 47  
}
```

Response (At Limit):

```
json  
  
{  
  "allowed": false,  
  "currentUsage": 100  
}
```

Use Case: Gateway wants to check before forwarding request.

3. Record Request

Purpose: Record that a request was made (increments counter)

Request:

```
http  
POST /api/rate-limit/record?apiKey=a1b2c3d4e5f6g7h8i9j0
```

Response (Success):

```
json  
{  
  "success": true,  
  "currentUsage": 48  
}
```

Response (Exceeded):

```
json  
HTTP 429 Too Many Requests  
{  
  "error": "Rate limit exceeded"  
}
```

Use Case: After successfully processing a request, record it.

4. Get Current Usage

Purpose: Check how many requests consumer has made in current window

Request:

```
http  
GET /api/rate-limit/usage?apiKey=a1b2c3d4e5f6g7h8i9j0&windowType=MINUTE
```

Response:

```
json
```

```
{  
  "apiKey": "a1b2c3d4e5f6g7h8i9j0",  
  "windowType": "MINUTE",  
  "currentUsage": 47  
}
```

Use Case: Dashboard showing real-time usage statistics.

5. Suspend Consumer

Purpose: Block a consumer from making requests (abuse, non-payment, etc.)

Request:

```
http  
PATCH /api/consumers/1/suspend
```

Response:

```
204 No Content
```

Effect: All future requests with this consumer's API key will be rejected with **403 Forbidden**.

6. Activate Consumer

Purpose: Re-enable a suspended consumer

Request:

```
http  
PATCH /api/consumers/1/activate
```

Response:

```
204 No Content
```

Real-World Scenario

Scenario: Weather API Service

You're running a weather API. You have 3 customers:

Customers:

1. Free Tier App: 10 requests/minute
2. Pro App: 100 requests/minute
3. Enterprise App: 1000 requests/minute

Timeline of Events

10:30:00 - Free Tier App starts hammering the API

```
10:30:05 → Request #1 → ✓ Allowed (1/10)
10:30:10 → Request #2 → ✓ Allowed (2/10)
10:30:15 → Request #3 → ✓ Allowed (3/10)
...
10:30:50 → Request #10 → ✓ Allowed (10/10)
10:30:55 → Request #11 → ✗ BLOCKED (429 Too Many Requests)
```

10:31:00 - New minute, counter resets

```
10:31:05 → Request #1 (of new window) → ✓ Allowed (1/10)
```

10:31:30 - Pro App makes a batch request

Sends 50 requests simultaneously

All arrive at same time (10:31:30)

System handles:

Request 1 → Creates record with count=1

Request 2 → Updates record to count=2

Request 3 → Updates record to count=3

...

Request 50 → Updates record to count=50

All 50 requests: ✓ ALLOWED (under 100 limit)

10:32:00 - Admin detects Free Tier abuse

Admin reviews logs:

Free Tier App tried to make 50 requests in one minute

Only 10 were allowed, 40 were blocked

Admin decision: SUSPEND

Action:

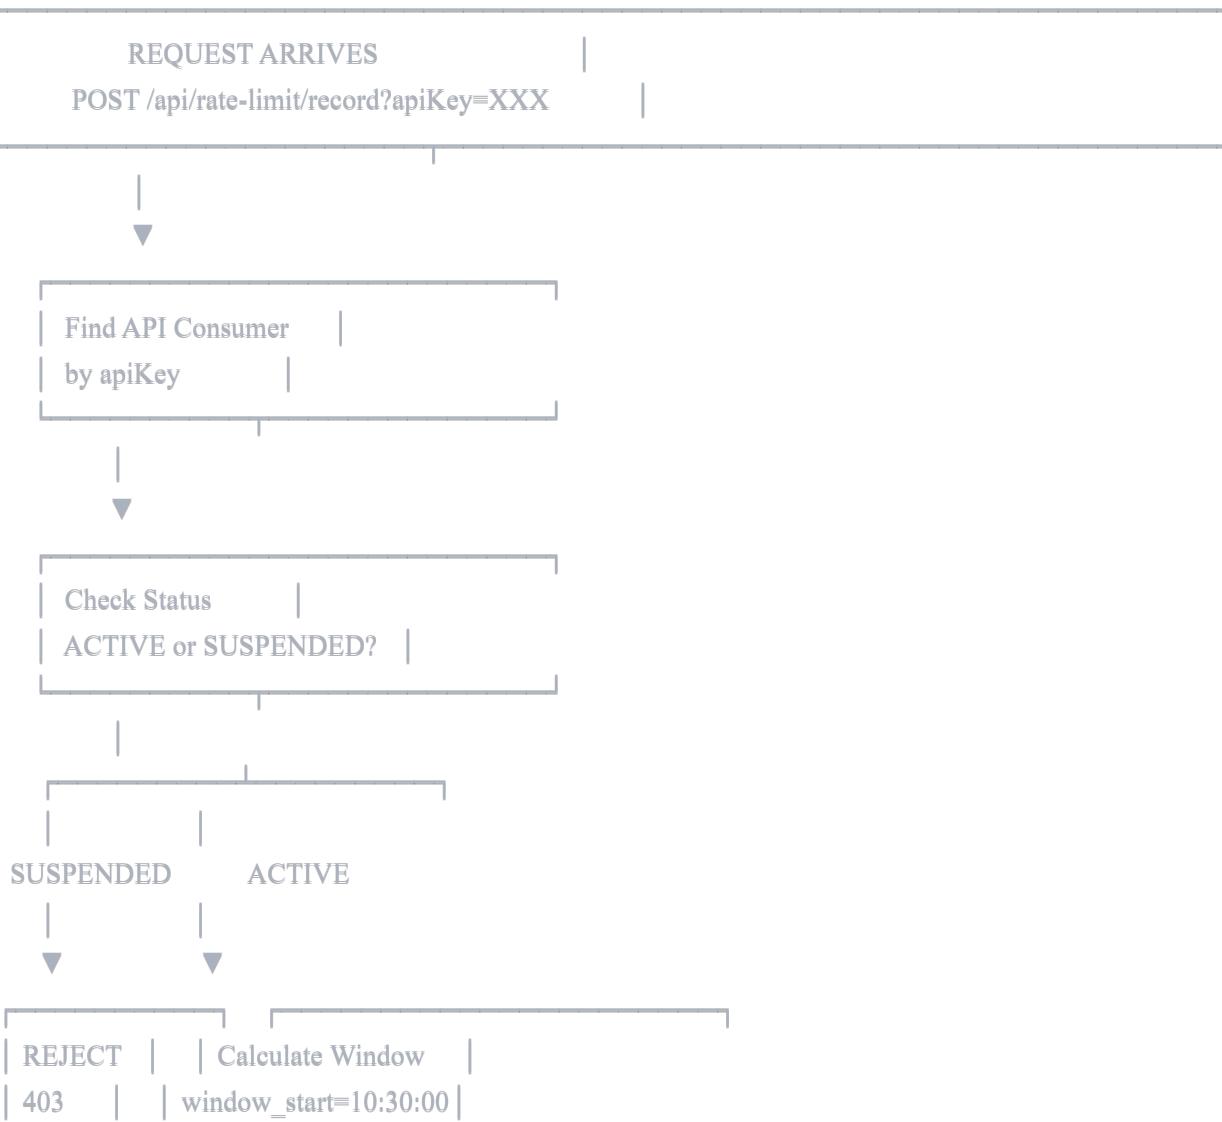
PATCH /api/consumers/1/suspend

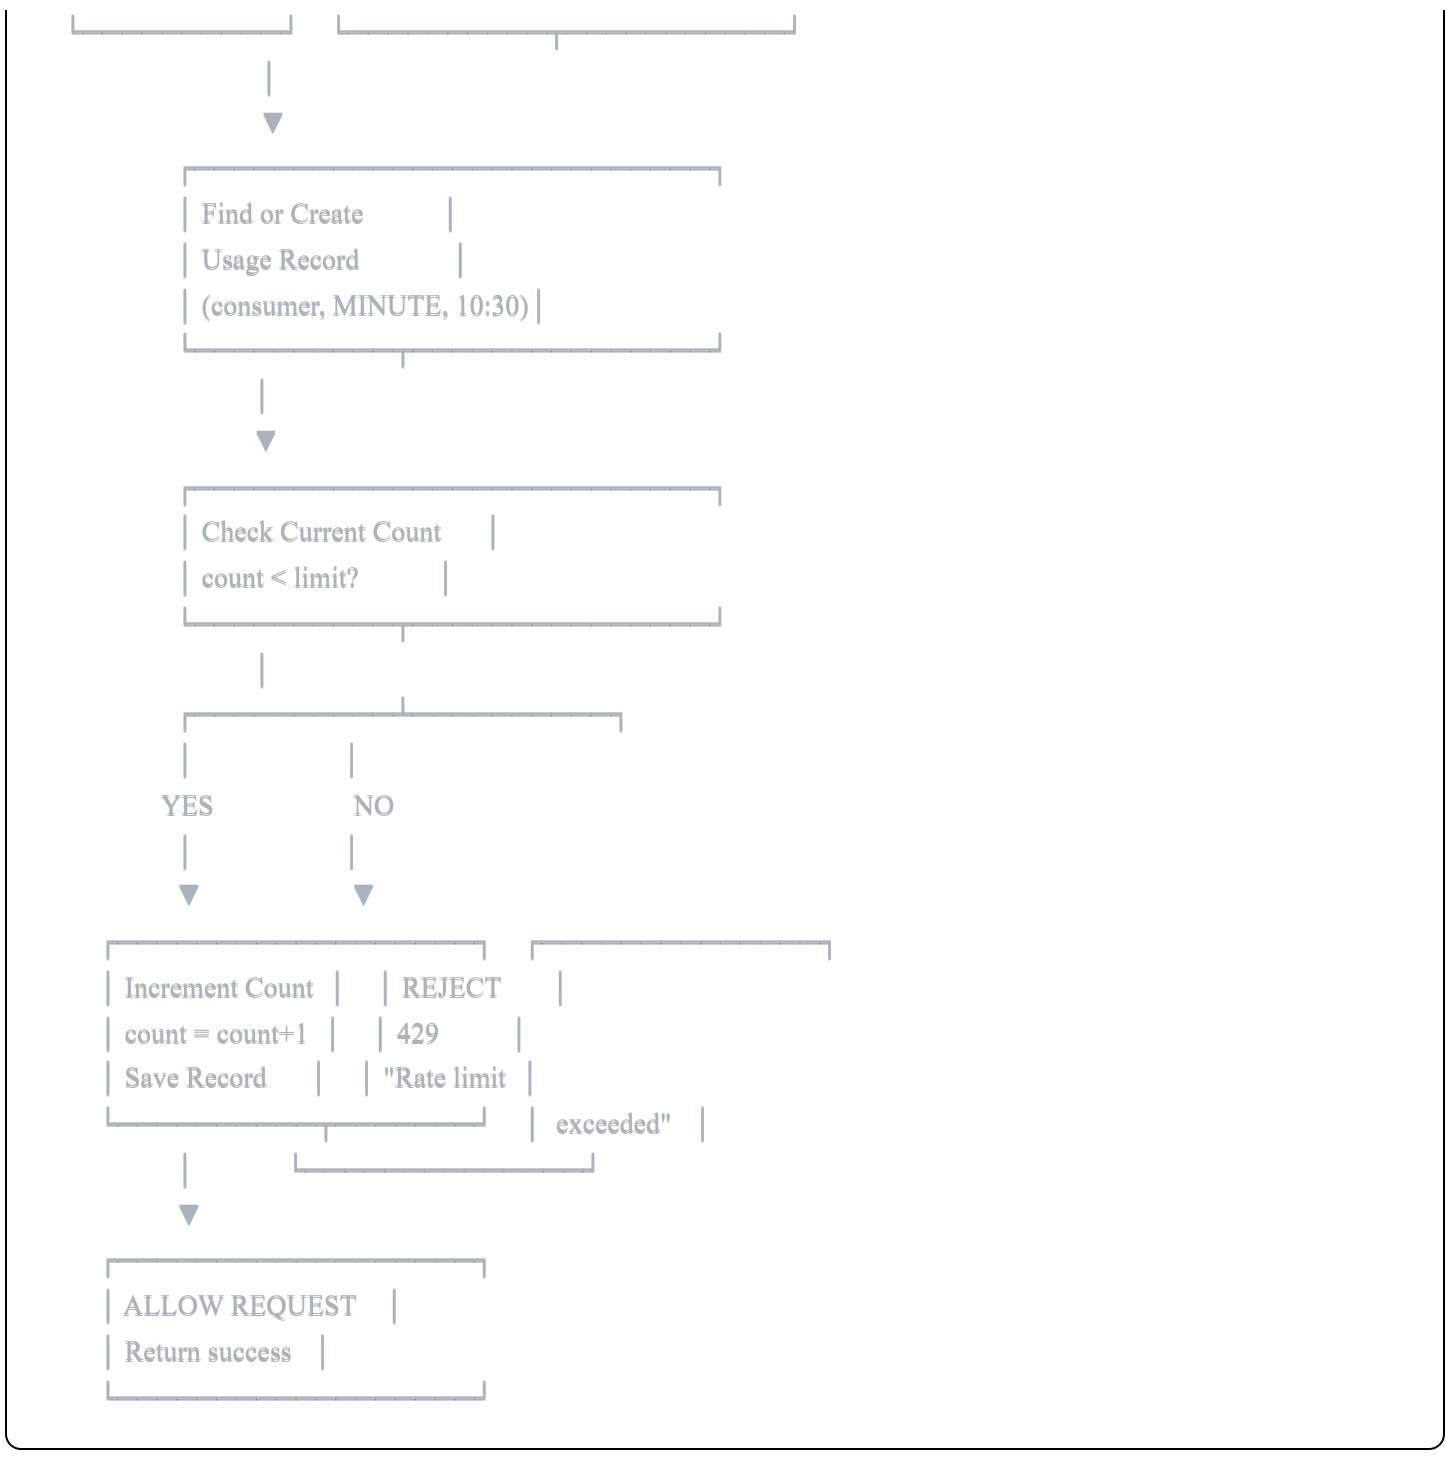
Result:

All future requests from Free Tier App → 403 Forbidden

"Consumer is suspended"

Summary: How Everything Connects





Key Takeaways

1. **API Consumer** = Identity (who)
2. **API Key** = Authentication (proof of identity)
3. **Rate Limit** = Rule (how many allowed)
4. **Time Window** = Period (when counter resets)
5. **Usage Record** = Counter (current usage)
6. **Unique Constraint** = Accuracy (one truth per window)

The system protects your API from overuse while providing fair access to all consumers.